


Principal Component Analysis (PCA)

Data Reduction

- summarization of data with many (p) variables by a smaller set of (k) derived (synthetic, composite) variables.



Data Reduction

- “Residual” variation is information in A that is not retained in X
- balancing act between
 - clarity of representation, ease of understanding
 - oversimplification: loss of important or relevant information.

Principal Component Analysis (PCA)

- probably the most widely-used and well-known of the “standard” multivariate methods
- invented by Pearson (1901) and Hotelling (1933)
- Also known under the name “factor analysis” (“principal factor analysis” is a synonym of PCA).

Principal Component Analysis (PCA)

- takes a data matrix of n objects by p variables, which may be correlated, and summarizes it by uncorrelated axes (principal components or principal axes) that are linear combinations of the original p variables
- the first k components display as much as possible of the variation among objects.

Geometric Rationale of PCA

- objects are represented as a cloud of n points in a multidimensional space with an axis for each of the p variables
- the **centroid** of the points is defined by the mean of each variable
- the **variance** of each variable is the average squared deviation of its n values around the mean of that variable.

$$V_i = \frac{1}{n-1} \sum_{m=1}^n (X_{im} - \bar{X}_i)^2$$

Geometric Rationale of PCA

- degree to which the variables are linearly correlated is represented by their **covariances**.

$$C_{ij} = \frac{1}{n-1} \sum_{m=1}^n (x_{im} - \bar{x}_i)(x_{jm} - \bar{x}_j)$$

Covariance of variables i and j

Sum over all n objects

Value of variable i in object m

Mean of variable i

Value of variable j in object m


Mean of variable j

Geometric Rationale of PCA

- objective of PCA is to rigidly rotate the axes of this p -dimensional space to new positions (principal axes) that have the following properties:
 - ordered such that principal axis 1 has the highest variance, axis 2 has the next highest variance, , and axis p has the lowest variance
 - covariance among each pair of the principal axes is zero (the principal axes are uncorrelated).

2D Example of PCA

- variables X_1 and X_2 have positive covariance & each has a similar variance.




$$V_1 = 6.67$$

$$V_2 = 6.24$$

$$C_{1,2} = 3.42$$


Configuration is Centered

- each variable is adjusted to a mean of zero (by subtracting the mean from each value).



Principal Components are Computed

- PC 1 has the highest possible variance (9.88)
- PC 2 has a variance of 3.03
- PC 1 and PC 2 have zero covariance.



The Dissimilarity Measure Used in PCA is Euclidean Distance

- PCA uses Euclidean Distance calculated from the p variables as the measure of dissimilarity among the n objects
- PCA derives the best possible k dimensional ($k < p$) representation of the Euclidean distances among objects.


Generalization to p -dimensions

- In practice nobody uses PCA with only 2 variables
- The algebra for finding principal axes readily generalizes to p variables
- PC 1 is the direction of maximum variance in the p -dimensional cloud of points
- PC 2 is in the direction of the next highest variance, subject to the constraint that it has zero covariance with PC 1.


Generalization to p -dimensions

- PC 3 is in the direction of the next highest variance, subject to the constraint that it has zero covariance with both PC 1 and PC 2
- and so on... up to PC p

- each principal axis is a linear combination of the original two variables
- $PC_j = a_{i1}Y_1 + a_{i2}Y_2 + \dots + a_{in}Y_n$
- a_{ij} 's are the coefficients for factor i , multiplied by the measured value for variable j



- PC axes are a rigid rotation of the original variables
- PC 1 is simultaneously the direction of maximum variance and a least-squares “line of best fit” (squared distances of points away from PC 1 are minimized).



Generalization to p -dimensions

- if we take the first k principal components, they define the k -dimensional “hyperplane of best fit” to the point cloud
- of the total variance of all p variables:
 - PCs 1 to k represent the maximum possible proportion of that variance that can be displayed in k dimensions
 - *i.e.* the squared Euclidean distances among points calculated from their coordinates on PCs 1 to k are the best possible representation of their squared Euclidean distances in the full p dimensions.

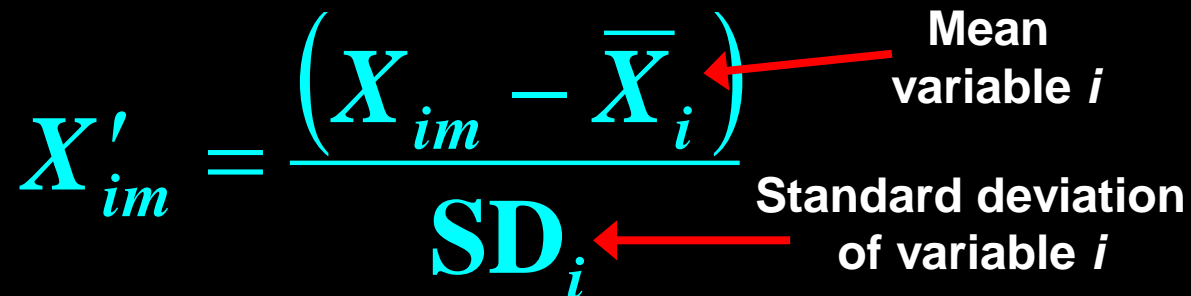
Covariance vs Correlation

- using covariances among variables only makes sense if they are measured in the same units
- even then, variables with high variances will dominate the principal components
- these problems are generally avoided by standardizing each variable to unit variance and zero mean.

$$X'_{im} = \frac{(X_{im} - \bar{X}_i)}{SD_i}$$

Mean variable i

Standard deviation of variable i



Covariance vs Correlation

- covariances between the standardized variables are **correlations**
- after standardization, each variable has a variance of 1.000
- correlations can be also calculated from the variances and covariances:

Correlation between variables i and j $\rightarrow r_{ij} = \frac{C_{ij}}{\sqrt{V_i V_j}}$

Covariance of variables i and j $\rightarrow C_{ij}$

Variance of variable i $\rightarrow V_i$

Variance of variable j $\rightarrow V_j$

The diagram illustrates the formula for calculating the correlation coefficient r_{ij} between two variables i and j . The formula is $r_{ij} = \frac{C_{ij}}{\sqrt{V_i V_j}}$. Red arrows point from the text labels to the corresponding parts of the formula: r_{ij} is labeled 'Correlation between variables i and j ', C_{ij} is labeled 'Covariance of variables i and j ', V_i is labeled 'Variance of variable i ', and V_j is labeled 'Variance of variable j '.

The Algebra of PCA

- first step is to calculate the **cross-products matrix** of variances and covariances (or correlations) among every pair of the p variables
- square, symmetric matrix
- diagonals are the variances, off-diagonals are the covariances.

	X_1	X_2
X_1	6.6707	3.4170
X_2	3.4170	6.2384

Variance-covariance Matrix

	X_1	X_2
X_1	1.0000	0.5297
X_2	0.5297	1.0000

Correlation Matrix

The Algebra of PCA

- in matrix notation, this is computed as

$$S = X'X$$

- where X is the $n \times p$ data matrix, with each variable centered (also standardized by SD if using correlations).

	X_1	X_2
X_1	6.6707	3.4170
X_2	3.4170	6.2384

Variance-covariance Matrix

	X_1	X_2
X_1	1.0000	0.5297
X_2	0.5297	1.0000

Correlation Matrix

Manipulating Matrices

- transposing: could change the columns to rows or the rows to columns

$$X = \begin{bmatrix} 10 & 0 & 4 \\ 7 & 1 & 2 \end{bmatrix}$$

$$X' = \begin{bmatrix} 10 & 7 \\ 0 & 1 \\ 4 & 2 \end{bmatrix}$$

- multiplying matrices
 - must have the same number of columns in the premultiplicand matrix as the number of rows in the postmultiplicand matrix

The Algebra of PCA

- sum of the diagonals of the variance-covariance matrix is called the **trace**
- it represents the **total variance** in the data
- it is the mean squared Euclidean distance between each object and the centroid in p -dimensional space.

	X_1	X_2
X_1	6.6707	3.4170
X_2	3.4170	6.2384

Trace = 12.9091

	X_1	X_2
X_1	1.0000	0.5297
X_2	0.5297	1.0000

Trace = 2.0000

The Algebra of PCA

- finding the principal axes involves eigenanalysis of the cross-products matrix (S)
- the eigenvalues (latent roots) of S are solutions (λ) to the characteristic equation

$$|S - \lambda I| = 0$$

The Algebra of PCA

- the eigenvalues, $\lambda_1, \lambda_2, \dots, \lambda_p$ are the variances of the coordinates on each principal component axis
- the sum of all p eigenvalues equals the trace of S (the sum of the variances of the original variables).

	X_1	X_2
X_1	6.6707	3.4170
X_2	3.4170	6.2384

$$\lambda_1 = 9.8783$$

$$\lambda_2 = 3.0308$$


$$\text{Trace} = 12.9091$$

$$\text{Note: } \lambda_1 + \lambda_2 = 12.9091$$

The Algebra of PCA

- each eigenvector consists of p values which represent the “contribution” of each variable to the principal component axis
- eigenvectors are uncorrelated (orthogonal)
 - their cross-products are zero.

Eigenvectors



	u_1	u_2
x_1	0.7291	-0.6844
x_2	0.6844	0.7291

$$0.7291 * (-0.6844) + 0.6844 * 0.7291 = 0$$

The Algebra of PCA

- coordinates of each object i on the k^{th} principal axis, known as the **scores** on PC k , are computed as

$$z_{ki} = u_{1k}x_{1i} + u_{2k}x_{2i} + \cdots + u_{pk}x_{pi}$$

- where Z is the $n \times k$ matrix of **PC scores**, X is the $n \times p$ **centered data matrix** and U is the $p \times k$ **matrix of eigenvectors**.


The Algebra of PCA

- variance of the scores on each PC axis is equal to the corresponding eigenvalue for that axis
- the eigenvalue represents the variance displayed (“explained” or “extracted”) by the k^{th} axis
- the sum of the first k eigenvalues is the variance explained by the k -dimensional ordination.

$\lambda_1 = 9.8783$ $\lambda_2 = 3.0308$ Trace = 12.9091

PC 1 displays ("explains")

$9.8783/12.9091 = 76.5\%$ of the total variance



The Algebra of PCA

- The cross-products matrix computed among the p principal axes has a simple form:
 - all off-diagonal values are zero (the principal axes are uncorrelated)
 - the diagonal values are the eigenvalues.

	PC_1	PC_2
PC_1	9.8783	0.0000
PC_2	0.0000	3.0308

Variance-covariance Matrix
of the PC axes

A more challenging example

- data from research on habitat definition in the endangered Baw Baw frog
- 16 environmental and structural variables measured at each of 124 sites
- correlation matrix used because variables have different units



Eigenvalues

Axis	Eigenvalue	% of Variance	Cumulative % of Variance
1	5.855	36.60	36.60
2	3.420	21.38	57.97
3	1.122	7.01	64.98
4	1.116	6.97	71.95
5	0.982	6.14	78.09
6	0.725	4.53	82.62
7	0.563	3.52	86.14
8	0.529	3.31	89.45
9	0.476	2.98	92.42
10	0.375	2.35	94.77

Interpreting Eigenvectors


- correlations between variables and the principal axes are known as **loadings**
- each element of the eigenvectors represents the contribution of a given variable to a component

	1	2	3
Altitude	0.3842	0.0659	-0.1177
pH	-0.1159	0.1696	-0.5578
Cond	-0.2729	-0.1200	0.3636
TempSurf	0.0538	-0.2800	0.2621
Relief	-0.0765	0.3855	-0.1462
maxERht	0.0248	0.4879	0.2426
avERht	0.0599	0.4568	0.2497
%ER	0.0789	0.4223	0.2278
%VEG	0.3305	-0.2087	-0.0276
%LIT	-0.3053	0.1226	0.1145
%LOG	-0.3144	0.0402	-0.1067
%W	-0.0886	-0.0654	-0.1171
H1Moss	0.1364	-0.1262	0.4761
DistSWH	-0.3787	0.0101	0.0042
DistSW	-0.3494	-0.1283	0.1166
DistMF	0.3899	0.0586	-0.0175

How many axes are needed?

- does the $(k+1)^{th}$ principal axis represent more variance than would be expected by chance?
- several tests and rules have been proposed
- a common “rule of thumb” when PCA is based on correlations is that axes with eigenvalues > 1 are worth interpreting

Baw Baw Frog - PCA of 16 Habitat Variables



What are the assumptions of PCA?

- assumes relationships among variables are LINEAR
 - cloud of points in p -dimensional space has linear dimensions that can be effectively summarized by the principal axes
- if the structure in the data is NONLINEAR (the cloud of points twists and curves its way through p -dimensional space), the principal axes will not be an efficient and informative summary of the data.